

# MLOps Pipeline Documentation


## Complete Guide for Cats vs Dogs Classification MLOps Pipeline


This document provides comprehensive documentation for the end-to-end MLOps pipeline implementation for binary image classification (Cats vs Dogs) for a pet adoption platform.

## Table of Contents

1. [Architecture Overview](#)
2. [Milestone 1: Model Development & Experiment Tracking](#)
3. [Milestone 2: Model Packaging & Containerization](#)
4. [Milestone 3: CI Pipeline](#)
5. [Milestone 4: CD Pipeline & Deployment](#)
6. [Milestone 5: Monitoring & Logging](#)
7. [Workflow Demonstration](#)

## Architecture Overview





## Milestone 1: Model Development & Experiment Tracking {#milestone-1}

### 1.1 Data & Code Versioning

#### Git for Source Code

```
# Initialize repository
git init
git add .
git commit -m "Initial commit"

# Create feature branch
git checkout -b feature/model-improvements
```

#### DVC for Data Versioning

```
# Initialize DVC
dvc init
```

```
# Add data to DVC tracking
dvc add data/raw
git add data/raw.dvc .gitignore
git commit -m "Add raw data tracking"

# Remote storage is configured to Dagshub
# URL:
https://dagshub.com/vishalvishal099/BinaryImageClassification_For_A_Pet_Adoption_Platform.dvc

# Push data to Dagshub
dvc push

# Pull data from Dagshub
dvc pull
```

## 1.2 Data Preprocessing

The preprocessing pipeline ([src/data/preprocess.py](#)) handles:

- **Image Loading:** Support for JPEG, PNG, and other formats
- **Resizing:** All images resized to 224×224 pixels
- **Color Space:** Convert to RGB (handles grayscale, RGBA)
- **Split:** 80% train / 10% validation / 10% test
- **Augmentation:** Horizontal flip, brightness/contrast, rotation, noise

```
# Run preprocessing
python src/data/preprocess.py --raw-dir data/raw --processed-dir
data/processed
```

## 1.3 Model Building

Implemented models in [src/models/cnn.py](#):

1. **SimpleCNN:** Baseline 4-layer CNN with batch normalization
2. **ImprovedCNN:** Enhanced CNN with residual-like connections
3. **ResNet18:** Transfer learning option with pretrained weights

```
from src.models.cnn import get_model

# Get baseline model
model = get_model("simple_cnn", num_classes=2, dropout_rate=0.5)

# Get improved model
model = get_model("improved_cnn", num_classes=2)
```

```
# Get pretrained ResNet
model = get_model("resnet18", pretrained=True)
```

## 1.4 Experiment Tracking with MLflow

```
# Train with MLflow tracking
python src/training/train.py --config configs/train_config.yaml

# Start MLflow tracking server (port 5001)
mlflow server --host 0.0.0.0 --port 5001 --backend-store-uri
sqlite:///mlflow.db
# UI: http://localhost:5001
# Dagshub remote:
https://dagshub.com/vishalvishal099/BinaryImageClassification_For_A_Pet_Ad
option_Platform.mlflow
```

### Logged Artifacts:

- Model parameters (batch\_size, learning\_rate, epochs, etc.)
- Metrics (train/val loss, accuracy, precision, recall, F1)
- Confusion matrix
- Trained model file
- Training curves

## Milestone 2: Model Packaging & Containerization {#milestone-2}

### 2.1 Inference Service (FastAPI)

The inference service (`src/inference/app.py`) provides:

Endpoint	Method	Description
/	GET	API information
/health	GET	Health check with model status
/predict	POST	Image classification
/metrics	GET	Prometheus metrics

```
# Example API response for /predict
{
    "predicted_class": "cat",
    "predicted_label": 0,
    "confidence": 0.95,
    "probabilities": {
        "cat": 0.95,
        "dog": 0.05
    }
}
```

```
    },
    "processing_time_ms": 45.2
}
```

## 2.2 Dependencies (requirements.txt)

All dependencies are pinned for reproducibility:

- `torch==2.1.0`
- `fastapi==0.104.1`
- `mlflow==2.8.1`
- `prometheus-client==0.19.0`

## 2.3 Containerization

```
# Build image
docker build -t cats-dogs-classifier:latest .

# Run container
docker run -p 8000:8000 -v ./models:/app/models cats-dogs-
classifier:latest

# Test with curl
curl http://localhost:8000/health

# Test prediction
curl -X POST -F "file=@test_image.jpg" http://localhost:8000/predict
```

---

# Milestone 3: CI Pipeline {#milestone-3}

## 3.1 Automated Testing

Tests are located in `tests/`:

```
# Run all tests
pytest tests/ -v

# Run with coverage
pytest tests/ --cov=src --cov-report=html
```

### Test Coverage:

- `test_preprocess.py`: Data preprocessing functions
- `test_inference.py`: API endpoints and inference logic

## 3.2 GitHub Actions CI Pipeline

The CI pipeline (`.github/workflows/ci.yml`) runs on every push/PR:

```
Jobs:  
1. lint      - Code quality checks (black, isort, flake8)  
2. test      - Unit tests with pytest  
3. build     - Docker image build and push  
4. integration - Docker container integration tests  
5. security   - Trivy vulnerability scanning
```

### 3.3 Artifact Publishing

Docker images are pushed to GitHub Container Registry:

```
ghcr.io/your-username/cats-dogs-classifier:latest  
ghcr.io/your-username/cats-dogs-classifier:<sha>
```

---

## Milestone 4: CD Pipeline & Deployment {#milestone-4}

### 4.1 Kubernetes Deployment

Manifests in `k8s/`:

File	Purpose
<code>namespace.yaml</code>	Dedicated namespace
<code>configmap.yaml</code>	Environment configuration
<code>deployment.yaml</code>	Pod deployment with 2 replicas
<code>service.yaml</code>	LoadBalancer and ClusterIP services
<code>hpa.yaml</code>	Horizontal Pod Autoscaler

```
# Deploy to Kubernetes  
kubectl apply -f k8s/  
  
# Check deployment  
kubectl get all -n cats-dogs-classifier
```

### 4.2 Docker Compose (Local)

```
# Start all services  
docker-compose up -d
```

```
# With monitoring stack
docker-compose --profile monitoring up -d

# View logs
docker-compose logs -f classifier
```

## 4.3 GitOps with ArgoCD

```
# Install ArgoCD
kubectl create namespace argocd
kubectl apply -n argocd -f
https://raw.githubusercontent.com/argoproj/argo-
cd/stable/manifests/install.yaml

# Apply application
kubectl apply -f k8s/argocd-application.yaml
```

## 4.4 Smoke Tests

```
# Run smoke tests
python scripts/smoke_test.py --url http://localhost:8000 --wait

# Output:
#  Health Check: PASS
#  Root Endpoint: PASS
#  Metrics Endpoint: PASS
#  Prediction Endpoint: PASS
```

---

# Milestone 5: Monitoring & Logging {#milestone-5}

## 5.1 Structured Logging

Request/response logging with structlog:

```
{
  "event": "prediction_completed",
  "predicted_class": "cat",
  "confidence": 0.95,
  "processing_time_ms": 45.2,
  "timestamp": "2024-01-15T10:30:00Z"
}
```

## 5.2 Prometheus Metrics

Available metrics at [/metrics](#):

Metric	Type	Description
<code>prediction_requests_total</code>	Counter	Total predictions by class
<code>prediction_latency_seconds</code>	Histogram	Request latency distribution
<code>health_check_requests_total</code>	Counter	Health check requests
<code>prediction_errors_total</code>	Counter	Prediction errors by type

### 5.3 Model Performance Tracking

The `PerformanceTracker` class ([src/utils/performance\\_tracker.py](#)):

- Collects predictions with timestamps
- Accepts ground truth labels for feedback
- Computes accuracy, precision, recall, F1
- Tracks latency statistics (mean, p95, p99)
- Persists metrics to JSON files

```
from src.utils.performance_tracker import get_tracker

tracker = get_tracker()
tracker.log_prediction(
    prediction="cat",
    confidence=0.95,
    latency_ms=45.2,
    ground_truth="cat" # Optional
)

stats = tracker.get_current_stats()
```

---

## Workflow Demonstration {#workflow-demonstration}

### Complete MLOps Workflow

```
# 1. Setup environment
make setup

# 2. Download and preprocess data
make download
make preprocess

# 3. Train model with MLflow tracking
make train

# 4. Run tests
```

```
make test

# 5. Build Docker image
make build

# 6. Run locally
make run

# 7. Test the API
make smoke-test

# 8. Deploy to Kubernetes
make deploy-k8s

# 9. Monitor with Prometheus
docker-compose --profile monitoring up -d
```

## Code Change to Deployed Model

### 1. Developer makes code change

```
git checkout -b feature/improve-model
# ... make changes ...
git add .
git commit -m "Improve model accuracy"
git push origin feature/improve-model
```

### 2. CI Pipeline runs automatically

- Linting and code quality checks
- Unit tests with pytest
- Docker image build
- Integration tests
- Security scanning

### 3. Merge to main triggers CD

- New Docker image pushed to registry
- Kubernetes deployment updated
- ArgoCD syncs changes
- Smoke tests verify deployment

### 4. Monitor deployed service

- Prometheus collects metrics
- Grafana dashboards show performance
- Logs available for debugging

# Quick Reference

## Important Commands

```

# Setup
make setup           # Full setup
source venv/bin/activate

# Development
make lint            # Run linters
make test             # Run tests
make train             # Train model

# Docker
make build            # Build image
make run               # Run container
make compose-up        # Start docker-compose

# Kubernetes
make deploy-k8s       # Deploy to K8s
make smoke-test        # Run smoke tests

# DVC
dvc pull              # Pull data
dvc push              # Push data
dvc repro             # Run pipeline

```

## Environment Variables

Variable	Description	Default
MODEL_PATH	Path to model file	/app/models/best_model.pt
PORT	Service port	8000
LOG_LEVEL	Logging level	INFO

## Troubleshooting

### Common Issues

#### 1. Model not loaded

- Ensure model file exists at MODEL\_PATH
- Check model checkpoint format

#### 2. Docker build fails

- Verify requirements.txt is correct
- Check Dockerfile stages

### 3. Kubernetes deployment fails

- Check PVC is bound
- Verify image pull secrets

### 4. Tests failing

- Install test dependencies: `pip install pytest pytest-cov`
  - Run with verbose: `pytest -v`
-